

**Problem 1** (3 parts, 30 points)

**Compilation, Concurrency & Interrupts**

**Part A** (20 points) Perform at least **five** standard compiler optimizations on the following C code fragment by writing the optimized version (in C) to the right.

|  |               |  |
|--|---------------|--|
| <pre>int foo(int x, int y) {     int g = 32;      int h=0, a = 0, b = 1, i=100;      while ((x*y)&lt;i) {         a += g*(h+i);         b *= exp(a, (h+i));         i--;     }     return (b); }</pre> | $\Rightarrow$ | <pre>int foo(int x, int y){     int a=0,b=1,i=100;      int tmp = x*y;      while(tmp&lt;i){         a += i &lt;&lt; 5;         b *= exp(a, i);         i--;     }     return (b); }</pre> |
|--|---------------|--|

Briefly describe which standard compiler optimizations you applied:

1. **Dead code elimination** (declaration of g)
2. **constant propagation** (g=32 and h=0)
3. **algebraic simplification** (h+i = 0+i = i)
4. **strength reduction** (32\*i = i<<5)
5. **loop invariant removal** (x\*y)
6. **(also common subexpression elimination: (h+i) can be computed once and simplified)**

**Part B** (6 points) In the table below, draw lines to match the type of concurrency with the architectural support required for it.

| Type of concurrency                 | Draw lines here: | Architectural Support                      |
|-------------------------------------|------------------|--|
| Data-level parallelism (DLP)        |                  | Multicore system with shared memory        |
| Instruction-level parallelism (ILP) |                  | Multimedia ISA extensions (e.g., MMX, SSE) |
| Thread-level parallelism (TLP)      |                  | Pipelining                                 |

**Part C** (4 points) What is the difference between a nonmaskable interrupt (NMI) and an asynchronous interrupt?

**An NMI always has priority over asynchronous interrupts. It cannot be masked or overridden by a different interrupt. An asynchronous interrupt (e.g., from keystrokes) can be masked by another interrupt (e.g., from the graphics card).**

**Problem 2** (4 parts, 35 points)**Associative Sets**

Consider a hash table that is implemented using the following struct definitions.

```
typedef struct Entry {
    int      Key;
    int      Value;
    struct Entry *Next;
} Entry;

typedef struct {
    Entry    *Buckets[13];
    int      Size;
} HashTable;
```

**Part A** (5 points) If a hash table of type `HashTable` is created which contains **273** entries (each of type `Entry`), how many total words of storage are used by this hashtable? (Show work.)

$$273 * 3 + 13 + 1 = 833$$

**Part B** (5 points) Suppose the same `Values` that are contained in the 273 entries are stored in an indexed set (an array), each `Value` indexed by the entry's `Key`. Assume the `Key` can be any unsigned integer representable by 32 bits. How many total words of storage are used by the indexed set? (Show work.)

$$2^{32} \text{ entries} = 4\text{B words}$$

**Part C** (10 points) Suppose the bucket lists in the `Buckets` array in **Part A** contain a number of `Entry` objects, evenly distributed across the hash table buckets. Assume that *computing the hash function* takes an average of **five operations** and *comparing two keys* takes an average of **four operations**. Ignore effects of spatial and temporal reference locality. Suppose that **75%** of keys looked up are found in the hash table and **25%** are not found. How many of these operations would be required for the average lookup in the hash table described above if the bucket list is unsorted versus sorted? (Show work.)

number of operations when each bucket list is *unsorted*:  $5 + (3/4)(22/2)4 + (1/4)(21)4 = 5 + 3(11) + 21 = 59$

number of operations when each bucket list is *sorted*:  $5 + 4(21+1)/2 = 49$

**Part D** (15 points) Write a C subroutine called `CreateEntry` that takes two integers (named `K` and `V`), allocates an object of type `Entry` on the heap with fields `Key=K`, `Value=V`, and `Next=NULL`, and returns a pointer to the `Entry` object allocated. The subroutine should be sure to check whether there was enough space on the heap for the object to be allocated and if not, it should print an error message and exit.

```
Entry *CreateEntry(int K, int V) {
    Entry *NewE = (Entry *)malloc(sizeof(Entry));
    if (NewE == NULL) {
        printf("Insufficient memory");
        exit(1);
    }
    NewE->Key = K;
    NewE->Value = V;
    NewE->Next = NULL;
    return (NewE);
}
```

**Problem 3** (1 parts, 25 points)**Heap and Hash Table**

Consider an open hash table composed of a four-bucket table, with each bucket containing a variable length list. Each list entry has three slots <key, value, next> corresponding to the three word groupings in the entries section. The hash function is  $\text{key} \bmod \text{four}$ . Inserted entries are *appended to the end* of a bucket list. Deallocated entries are maintained on a LIFO free list. When the free list is empty, new entry objects are allocated from heap memory. Accesses are listed as <op, key, [value]>. Simulate the access list below and draw the ending state. Assume the hash table is initially empty, the heap pointer is initially 5016 and the free pointer is initially 0.

|                 |  |              |                                |
|-----------------|--|--------------|--------------------------------|
| Heap<br>Pointer | <del>5016</del> 5028 <del>5040</del> 5052 5064 | Free<br>List | 0000 <del>5016</del> 5052 5016 |
|-----------------|--|--------------|--------------------------------|

## Buckets

|      |  |      |      |      |      |      |                      |
|------|--|------|------|------|------|------|----------------------|
| 5000 |  | 5004 | 5028 | 5008 | 5052 | 5012 | <del>5016</del> 5040 |
|------|--|------|------|------|------|------|----------------------|

## Entries

|      |                     |      |                      |      |  |      |  |
|------|---------------------|------|----------------------|------|--|------|--|
| 5016 | 2003                | 5040 | 2007                 | 5064 |  | 5088 |  |
| 5020 | 111                 | 5044 | <del>333</del> 555   | 5068 |  | 5092 |  |
| 5024 | <del>0 5040</del> 0 | 5048 | 0                    | 5072 |  | 5096 |  |
| 5028 | 2001                | 5052 | <del>2005</del> 2002 | 5076 |  | 5100 |  |
| 5032 | 222                 | 5056 | <del>444</del> 666   | 5080 |  | 5104 |  |
| 5036 | <del>0 5052</del> 0 | 5060 | <del>0 5016</del> 0  | 5084 |  | 5108 |  |

## Hash Table Access Trace

| # | op     | key  | value | # | op     | key  | value |
|---|--------|------|-------|---|--------|------|-------|
| 1 | insert | 2003 | 111   | 5 | insert | 2007 | 555   |
| 2 | insert | 2001 | 222   | 6 | remove | 2003 | n/a   |
| 3 | insert | 2007 | 333   | 7 | remove | 2005 | n/a   |
| 4 | insert | 2005 | 444   | 8 | insert | 2002 | 666   |

**Problem 4** (4 parts, 30 points)**Garbage Collection, Function Pointers, and MIPS**

Below is a snapshot of heap storage. Values that are pointers are denoted with a “\$”. The heap pointer is \$6168. The heap has been allocated contiguously beginning at \$6000, with no gaps between objects.

| addr | value  | addr | value  | addr | value  | addr | value  | addr | value  | addr | value |
|------|--------|------|--------|------|--------|------|--------|------|--------|------|-------|
| 6000 | 16     | 6032 | 12     | 6064 | 0      | 6096 | 16     | 6128 | 12     | 6160 | 0     |
| 6004 | 33     | 6036 | 28     | 6068 | 4      | 6100 | 6172   | 6132 | 72     | 6164 | 0     |
| 6008 | \$6132 | 6040 | \$6120 | 6072 | \$6100 | 6104 | 16     | 6136 | 9      | 6168 | 0     |
| 6012 | 16     | 6044 | \$6080 | 6076 | 8      | 6108 | 5      | 6140 | \$6004 | 6172 | 0     |
| 6016 | 80     | 6048 | 16     | 6080 | \$6036 | 6112 | \$6148 | 6144 | 20     | 6176 | 0     |
| 6020 | 8      | 6052 | \$6092 | 6084 | 6012   | 6116 | 8      | 6148 | 6046   | 6180 | 0     |
| 6024 | 25     | 6056 | \$6024 | 6088 | 4      | 6120 | 32     | 6152 | 8      | 6184 | 0     |
| 6028 | \$6036 | 6060 | 0      | 6092 | \$6080 | 6124 | \$6024 | 6156 | 26     | 6188 | 0     |

**Part A** (6 points) Suppose the stack holds a local variable whose value is the memory address \$6052 and register \$3 holds the address \$6004. No other registers or static variables currently hold heap memory addresses. List the addresses of all objects in the heap that are *not* garbage.

Addresses of **6052, 6092, 6024, 6080, 6036, 6120, 6004, 6132**

Non-Garbage Objects: \_\_\_\_\_

**Part B** (6 points) If the local variable whose value is the address \$6052 is popped from the stack, which addresses will be reclaimed by each of the following strategies? If none, write “none.”

|                                 |   |
|---------------------------------|---|
| <b>Reference Counting:</b>      | <b>6052, 6092</b>                         |
| <b>Mark and Sweep:</b>          | <b>6052, 6092, 6024, 6036, 6120, 6080</b> |
| <b>Old-New Space (copying):</b> | <b>6052, 6092, 6024, 6036, 6120, 6080</b> |

**Part C** (8 points) Complete the C code below by following these two steps:

1. Create a local variable, called `compare`, in `My_Search` that is a function pointer that points to `GT` if `ascending` is nonzero and to `LT` otherwise. Define the function pointer type with `typedef`.
2. Pass this function pointer to the subroutine `Climb` as its third parameter.

```
int GT(int x, int y)
    return(x>y);
int LT(int x, int y)
    return(abs(x)<abs(y));

typedef _____ int (* FP) (int, int); _____; /* part 1*/
int My_Search(int a, int b, int ascending) {

    _____ FP compare; _____; /* part 1*/

    _____ compare = (ascending? GT : LT) _____; /* part 1*/

    Climb(a, b, _____ compare _____);          /* part 2 */
    ...rest of My_Search's body...
}
```

**Part D** (10 points) Write the MIPS code implementation of the following C program fragment.

```
int A[100] = {4, -1, 3, ..., 17};
int B[25];
int i;
for(i=0; i<25; i++)
    B[i] = A[4*i];
```

Modify only registers \$1, \$2, \$3, and \$4. *For maximum credit, include comments.*

| Label               | Instruction                          | Comment                                      |
|---------------------|--------------------------------------|--|
|                     | <code>.data</code>                   |  |
| <code>Aaddr:</code> | <code>.word 4, -1, 3, ..., 17</code> | <code># int A[100]={4,-1,3,...,17};</code>   |
| <code>Baddr:</code> | <code>.alloc 25</code>               | <code># int B[25];</code>                    |
|                     | <code>.text</code>                   |  |
|                     | <code>addi \$1, \$0, 0</code>        | <code># initialize loop counter</code>       |
| <code>Loop:</code>  | <code>slti \$2, \$1, 100</code>      | <code># is counter &lt; 100 (i&lt;25)</code> |
|                     | <code>beq \$2, \$0, Exit</code>      | <code># if not, Exit Loop</code>             |
|                     | <code>sll \$3, \$1, 2</code>         | <code># 4*i</code>                           |
|                     | <code>lw \$2, Aaddr(\$3)</code>      | <code># read A[4*i]</code>                   |
|                     | <code>sw \$2, Baddr(\$1)</code>      | <code># write it to B[i]</code>              |
|                     | <code>addi \$1, \$1, 4</code>        | <code># increment counter</code>             |
|                     | <code>j Loop</code>                  | <code># loop back</code>                     |
| <code>Exit:</code>  | <code>...</code>                     | <code># instructions after the loop</code>   |

**Problem 5** (1 parts, 30 points)

## Activation Frames

The function `Bar` (below left) calls function `Foo` after completing code block 1. Write MIPS assembly code that properly calls `Foo`. Include all instructions between code block 1 and code block 2. Symbolically label all required stack entries and give their values if they are known (below right).

```
int Bar() {
    int A[] = {5, 25};
    int N = 0;

    (code block 1)

    A[0] = Foo(A, A[1], &N)

    (code block 2)
}
```

|              |      |               |             |
|--------------|------|---------------|-------------|
| Bar's FP     | 9900 | XXX           | XXX         |
|              | 9896 | A[1]          | 25          |
|              | 9892 | A[0]          | 5           |
| SP           | 9888 | N             | 0           |
|              | 9884 | <b>FP</b>     | <b>9900</b> |
|              | 9880 | <b>RA</b>     | <b>N/A</b>  |
|              | 9876 | <b>A</b>      | <b>9892</b> |
|              | 9872 | <b>A[1]</b>   | <b>25</b>   |
|              | 9868 | <b>&amp;N</b> | <b>9888</b> |
| Foo's FP, SP | 9864 | <b>RV</b>     | <b>N/A</b>  |
|              | 9860 |               |             |
|              | 9856 |               |             |

| label | instruction                 | comment                       |
|-------|-----------------------------|-------------------------------|
|       | <b>addi \$29, \$29, -24</b> | # allocate activation frame   |
|       | <b>sw \$30, 20(\$29)</b>    | # preserve bookkeeping info   |
|       | <b>sw \$31, 16(\$29)</b>    |                               |
|       | <b>addi \$1, \$30, -8</b>   | # push inputs                 |
|       | <b>sw \$1, 12(\$29)</b>     |                               |
|       | <b>lw \$1, -4(\$30)</b>     |                               |
|       | <b>sw \$1, 8(\$29)</b>      |                               |
|       | <b>addi \$1, \$30, -12</b>  |                               |
|       | <b>sw \$1, 4(\$29)</b>      |                               |
|       | <b>jal Foo</b>              | # call Foo                    |
|       | <b>lw \$31, 16(\$29)</b>    | # restore bookkeeping info    |
|       | <b>lw \$30, 20(\$29)</b>    |                               |
|       | <b>lw \$1, 0(\$29)</b>      | # read return value           |
|       | <b>sw \$1, -8(\$30)</b>     | # store return value in A[0]  |
|       | <b>addi \$29, \$29, 24</b>  | # deallocate activation frame |